

# USER CHURN PROJECT | Logistic Regression Baseline Model



Prepared for the Waze Leadership Team

## Overview

The Waze data team is building predictive models to boost growth by identifying users at risk of churn. This part of the analysis focused on developing and evaluating a baseline **logistic regression model** as a simple, interpretable benchmark for comparison with more advanced models in later phases.

**This report summarizes Milestone 5 and its implications for future project development.**

## Project Status

### Milestone 5 - Regression Modeling


**Target Goal:** Fit a logistic regression model to predict churn and evaluate baseline performance.

#### Methods:

- **Cleaned and engineered features**, including ratio metrics and the professional\_driver flag.
- **Addressed multicollinearity** by removing highly correlated predictors.
- **Encoded categorical variables** (label\_2, device\_binary).
- **Standardized numeric features** and built a logistic regression pipeline.
- **Applied class weighting** to account for imbalance (82% retained, 18% churned).

**Impact:** Produced a baseline model capable of identifying churners at a much higher rate, directly **supporting Waze's business goal of churn prevention**.

- **Balanced performance:** The F1 score (0.43) highlights a reasonable balance between precision and recall, providing a useful baseline for intervention strategies.
- **Model strength:** A **ROC AUC of 0.74** confirms the model is meaningfully distinguishing churners from retained users beyond random guessing.



## Key Insights

- **Churn detection:** The baseline logistic regression model correctly identified **70% of churners**, giving the team strong coverage of at-risk users.
- **Trade-off in accuracy:** Overall accuracy was **67%**, reflecting the model's willingness to predict churn more often in order to avoid missing true churners.
- **Precision vs. recall:** When the model predicts churn, ~31% of those users actually churn. While this introduces false positives, it ensures the model captures the majority of true churners.

## Next Steps

- **Threshold tuning:** Adjust churn probability cutoff to explore different precision/recall trade-offs based on business needs.
- **Feature interpretation:** Examine model coefficients to highlight which behaviors are most predictive of churn.
- **Model comparisons:** Benchmark against more complex models to test for performance improvements.
- **Stakeholder alignment:** Collaborate with product and marketing to decide whether catching more churners (high recall) or reducing false alarms (higher precision) better supports Waze's retention strategy.